

Build a Web AI App

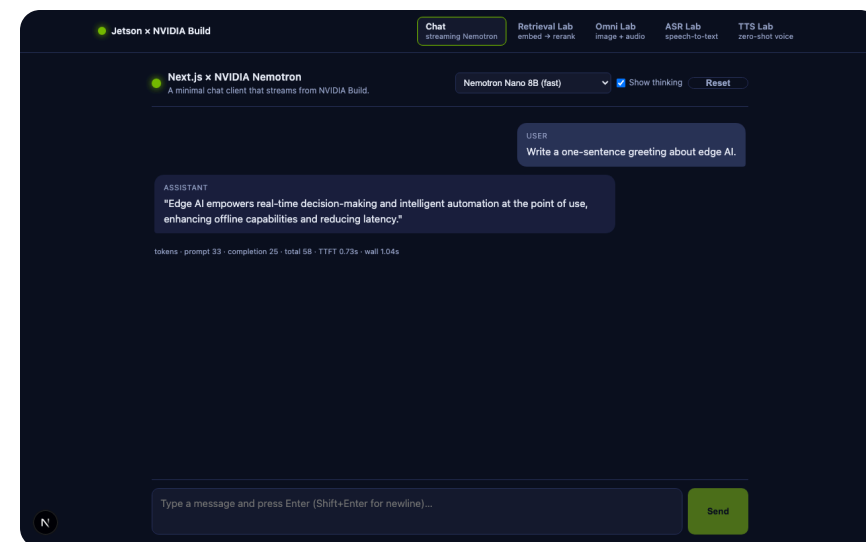
Next.js + NVIDIA Nemotron (Build API)

SJSU · Edge AI

A streaming chatbot web app that runs on your Jetson and serves to any browser.

1 What you'll build

- A **Next.js** web app with a **streaming chat** UI.
- Talks to **NVIDIA Nemotron** (and optionally **OpenAI** / **Anthropic**) via cloud APIs.
- **Runs on the Jetson**, opened from your laptop's browser.
- 🗝️ Your API key stays **server-side** — never sent to the browser.

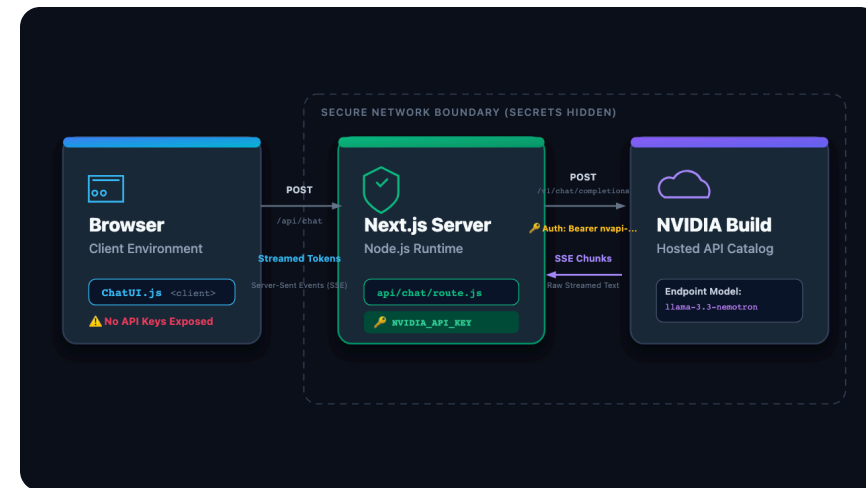


2 The big picture

Your browser talks to a small **server on the Jetson**, which talks to the model API:

Browser → /api/chat (on Jetson) →
NVIDIA/OpenAI/Anthropic → stream back

- The **page** runs in the browser (buttons, live text).
- The **API route** runs on the server — it holds the key and calls the model.
- The browser **never sees the key**.



3 Key web concepts (background)

- **Next.js** = a React framework. Each folder under `app/` is a page (the *App Router*).
- **Server Components / routes** run on the Jetson — they can hold **secrets** and call APIs.
- **Client Components** run in the **browser** — they're interactive (`"use client"`).
- **Streaming (SSE)**: the server forwards tokens one chunk at a time, so the answer appears live.

Rule of thumb: **secrets + external API calls** → **server**; **buttons + live updates** → **client**.

4 Where's the code (and what each part does)

```
edgeLLM/nextjs-nemotron-app/  
  app/page.js           # the chat UI - Client Component (browser)  
  app/components/ChatUI.js # streaming chat box  
  app/api/chat/route.js  # server route - holds the key, streams from the model  
  app/api/models/route.js # the model dropdown list  
  lib/providers.js       # picks NVIDIA/OpenAI/Anthropic + reads keys from ~/.env.local  
  .env.local            # optional local keys (git-ignored)
```

 Repo: [edgeLLM/nextjs-nemotron-app](https://github.com/edgeLLM/nextjs-nemotron-app)

5 Setup — API keys

Keys come from your `~/.env.local` (the same file `sjsujetsontool chat` saved). Add any of:

```
echo "NVIDIA_API_KEY=nvapi-..." >> ~/.env.local # build.nvidia.com (free)
echo "OPENAI_API_KEY=sk-..." >> ~/.env.local # platform.openai.com
echo "ANTHROPIC_API_KEY=sk-ant-..." >> ~/.env.local # console.anthropic.com
```

The app picks the provider from the **model you choose** (`nvidia/...` , `gpt-...` , `claude-...`).

 **No Node on the host or in the container yet?** That's normal — the next slide installs everything with one command. No `sudo` is needed anywhere.

6 Run it — `sjsujetsontool node`

One command installs Node in the container, runs `npm install`, and starts the dev server.
Run from **anywhere** on the host (your home is fine):

```
sjsujetsontool node          # interactive: prompts for path + mode
```

It asks the path with a sensible default — press *Enter* for this lesson:

```
● node v20.20.2 · npm 10.8.2 (inside container jetson-dev)
📁 Project path? [Enter = /Developer/edgeAI/edgeLLM/nextjs-nemotron-app]:
📁 Project: /Developer/edgeAI/edgeLLM/nextjs-nemotron-app
▶ Start the frontend now? [f]oreground / [b]ackground / [n]o:  b
🚀 Starting in BACKGROUND on port 3000.    • URL: http://192.168.5.206:3000
```

Shortcuts (skip the prompts):

```
sjsujetsontool node bg          # bg, default path
sjsujetsontool node fg /Developer/my-vite-app  # fg, explicit path (any order)
sjsujetsontool node /Developer/my-app bg      # path + mode, swapped
sjsujetsontool node stop        # stop a background server
```

Default path: `/Developer/edgeAI/edgeLLM/nextjs-nemotron-app` — override with
`SJSUJETSONTOOL_NODE_DIR=/Developer/foo` in your shell rc.

7 Manual install (what `sjsujetsontool node` does for you)

If you ever need to install Node by hand — or you want to see what the one-step command runs — open a container shell and use NodeSource's apt repo (Ubuntu 24.04 aarch64, root inside, *no sudo*):

```
sjsujetsontool shell          # drops into root@jetson-dev:/workspace
curl -fsSL https://deb.nodesource.com/setup_20.x | bash -
apt-get install -y nodejs     # → node v20.20.2 · npm 10.8.2
```

Then build and run the app the classic way:

```
cd /Developer/edgeAI/edgeLLM/nextjs-nemotron-app
npm install          # first time only (~30-60 s)
npm run dev          # serves on 0.0.0.0:3000
```

The Node install lives in the **container's writable layer** — it persists across `sjsujetsontool shell` invocations until the image is rebuilt. `node_modules/` lives on the host SSD because the parent dir is a host mount.

8 Open it from your laptop

Find the Jetson's IP and open it in any browser:

```
hostname -I | awk '{print $1}'          # e.g. 192.168.5.206 → http://192.168.5.206:3000
```

Each message streams **Jetson** → **model API** → **Jetson** → **your browser**, with a live TTFT and tokens-per-second line under the chat. To stop a backgrounded server:

```
sj sujetson tool node stop
```

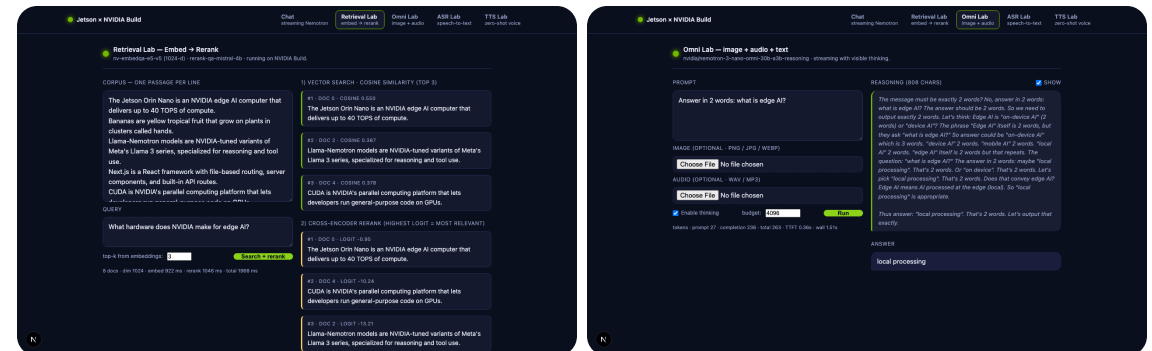
9 Extend it — same pattern every time

Every feature = **one page** (UI) + **one API route** (server logic). Copy the pattern:

```
app/<feature>/page.js      # the page/UI
app/api/<feature>/route.js  # server: read key, call a model, return/stream
```

The bonus labs are exactly this — add a route + page and you've extended the app:

 **Retrieval** ·  **Omni** (vision) ·  **ASR** ·  **TTS**



10 Make it your own (push to GitHub)

Copy the app into your own folder, then create your repo:

```
cp -r /Developer/edgeAI/edgeLLM/nextjs-nemotron-app ~/my-ai-app
cd ~/my-ai-app && rm -rf node_modules .next
git init && git add -A && git commit -m "My Edge AI web app"
```

Create an empty repo on github.com (the + → **New repository**), then:

```
git remote add origin https://github.com/<your-username>/my-ai-app.git
git branch -M main && git push -u origin main
```

With the GitHub CLI it's one line: `gh repo create my-ai-app --public --source=. --push`

Full walkthrough

Step-by-step build, code, and the four bonus labs:

lkk688.github.io/edgeAI/curriculum/11_nextjs_nemotron_app

← Back to the [Get-Started slides](#)